

Neural Machine Translation with Attention

January 31, 2024

[39]: `### v1.1`

```
[40]: from tensorflow.keras.layers import Bidirectional, Concatenate, Permute, Dot,   
      ↪Input, LSTM, Multiply   
      from tensorflow.keras.layers import RepeatVector, Dense, Activation, Lambda   
      from tensorflow.keras.optimizers import Adam   
      from tensorflow.keras.utils import to_categorical   
      from tensorflow.keras.models import load_model, Model   
      import tensorflow.keras.backend as K   
      import tensorflow as tf   
      import numpy as np   
   
      from faker import Faker   
      import random   
      from tqdm import tqdm   
      from babel.dates import format_date   
      from nmt_utils import *   
      import matplotlib.pyplot as plt   
      %matplotlib inline
```

```
[41]: # The model built here can be used to translate from one language to another.   
      ↪But Language Translation requires massive datasets   
      # and takes days of training on GPUs.   
      # So here we will perform a simpler 'date translation' task where the   
      # 'the 29th of August 1958' will be translated to '1958-08-29'   
      # '03/30/1968' will be translated to '1968-03-30'   
      # '24 JUNE 1987' will be translated to '1987-06-24'   
      # human readable date will be translated to machine readable date   
   
      m = 10000   
      dataset, human_vocab, machine_vocab, inv_machine_vocab = load_dataset(m)   
   
      # dataset is a list of tuples and each tuple is a pair of (input, output)   
      # 'human_vocab' is a dictionary mapping each character in   
      ↪human_readable_date(input) to an integer-valued index   
      # 'machine_vocab' is a dictionary mapping each character in   
      ↪machine_readable_date(output) to an integer-valued index
```

```
# There are 37 different characters in the 'human_vocab'. len(human_vocab) = 37
# There are 11 different characters in the 'machine_vocab'. len(machine_vocab)
→ = 11
```

```
100%|      | 10000/10000 [00:00<00:00, 24544.15it/s]
```

```
[42]: dataset[:10]
```

```
[42]: [('27 november 1980', '1980-11-27'),
      ('friday september 13 2019', '2019-09-13'),
      ('tuesday july 17 2018', '2018-07-17'),
      ('4/10/19', '2019-04-10'),
      ('wednesday april 27 1977', '1977-04-27'),
      ('tuesday december 6 1977', '1977-12-06'),
      ('01 sep 1991', '1991-09-01'),
      ('3 10 22', '2022-10-03'),
      ('tuesday july 20 1999', '1999-07-20'),
      ('wednesday january 29 1992', '1992-01-29')]
```

```
[43]: Tx = 30
      Ty = 10
      X, Y, Xoh, Yoh = preprocess_data(dataset, human_vocab, machine_vocab, Tx, Ty)

      print("X.shape:", X.shape)
      print("Y.shape:", Y.shape)
      print("Xoh.shape:", Xoh.shape)
      print("Yoh.shape:", Yoh.shape)

      # We set Tx = 30 which is the maximum length of the human readable date in our
      → dataset
      # We set Ty = 10 as the machine_readable_date is exactly 10 characters long

      # The following preprocessing is done

      # X: a processed version of the human readable dates in the training set.
      #   - Each character in X is replaced by an index (integer) mapped to the
      → character using human_vocab.
      #   - Each date is padded to ensure a length of      using a special character
      → (< pad >).
      #   - X.shape = (m, Tx) where m is the number of training examples in a batch.

      # Y: a processed version of the machine readable dates in the training set.
      #   - Each character is replaced by an index (integer) mapped to the character
      → using machine_vocab.
      #   - Y.shape = (m, Ty)
```

```
# Xoh: one-hot version of X
# Yoh: one-hot version of Y
```

```
X.shape: (10000, 30)
Y.shape: (10000, 10)
Xoh.shape: (10000, 30, 37)
Yoh.shape: (10000, 10, 11)
```

[44]: *# Let's look at an example of the preprocessed training example*

```
index = 0
print("Source date:", dataset[index][0])
print("Target date:", dataset[index][1])
print()
print("Source after preprocessing (indices):", X[index])
print("Target after preprocessing (indices):", Y[index])
print()
print("Source after preprocessing (one-hot):", Xoh[index])
print("Target after preprocessing (one-hot):", Yoh[index])
```

```
Source date: 27 november 1980
Target date: 1980-11-27
```

```
Source after preprocessing (indices): [ 5 10  0 25 26 32 17 24 14 17 28  0  4 12
11  3 36 36 36 36 36 36 36
36 36 36 36 36 36]
```

```
Target after preprocessing (indices): [ 2 10  9  1  0  2  2  0  3  8]
```

```
Source after preprocessing (one-hot): [[0. 0. 0. ... 0. 0. 0.]
[0. 0. 0. ... 0. 0. 0.]
[1. 0. 0. ... 0. 0. 0.]
...
[0. 0. 0. ... 0. 0. 1.]
[0. 0. 0. ... 0. 0. 1.]
[0. 0. 0. ... 0. 0. 1.]]
```

```
Target after preprocessing (one-hot): [[0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
[0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
[0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
[1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
[1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]
[0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0.]]
```

```
[45]: # Defined shared layers as global variables
# A vector of shape (m,n_s) becomes (m,Tx,n_s) after repeating
repeater = RepeatVector(Tx)

concatenator = Concatenate(axis=-1)
# concatenator(a,b) where a is of dimension (x,y,z1) and b is of dimension
→ (x,y,z2).
# All the dimensions need to match except the last dimension
# The resultant is of dimension (x,y,z1+z2)

# The 'Dense' layer only transforms the last dimension.
# When you pass a tensor of size (x,y,z) through the 'dense1', the resultant
→ is of size (x,y,10).
dense1 = Dense(10, activation = "tanh")
dense2 = Dense(1, activation = "relu")
# When you pass a tensor through a Dense layer, it multiplies the values in the
→ last dimension of the input tensor with
# the layer's weights and adds the bias, followed by applying the activation
→ function, resulting in a new last dimension as
# specified by the number of units in the Dense layer.

activator = Activation(softmax, name='attention_weights')
# We are using a custom softmax(axis = 1) loaded in this notebook

dotor = Dot(axes = 1)

# Note

# Question - The advantage of attention model is that it doesn't have to wait
→ till processing the entire sentence and can
# start the language translation even after few words of the input sentence are
→ processed right? but don't we take
# attention on x<1>, x<2>,....x<tx> all into consideration? Aren't we still
→ processing the complete input sentence and
# then only y<1> is being produced right?

# Answer
# Models that don't use attention put equal emphasis on every part of the input
→ sentence whereas models that
# use attention puts more emphasis on relevant parts of the input sentence.
→ This is the only difference.
# Otherwise both models with and without attention does depend on the entire
→ input sequence.
```

```
[46]: def one_step_attention(a, s_prev):
    """
    Performs one step of attention: Outputs a context vector computed as a dot_
    ↪product of the attention weights
    "alphas" and the hidden states "a" of the Bi-LSTM.

    Arguments:
    a -- hidden state output of the Bi-LSTM, numpy-array of shape (m, Tx, 2*n_a)
    s_prev -- previous hidden state of the (post-attention) LSTM, numpy-array_
    ↪of shape (m, n_s)

    Returns:
    context -- context vector, input of the next (post-attention) LSTM cell
    """

    # s_prev is of shape (m, n_s)
    # Use 'repeater' to repeat s_prev to be of shape (m, Tx, n_s) so that we_
    ↪can concatenate it with all hidden states "a"
    s_prev = repeater(s_prev)

    # Use 'concatenator' to concatenate a and s_prev on the last axis
    concat = concatenator([a,s_prev])
    # 'concat' is of shape (m, Tx, 2*n_a + n_s)

    # Use 'densor1' to propagate concat through a small fully-connected neural_
    ↪network to compute the
    # "intermediate energies" variable e.
    e = densor1(concat)
    # 'e' is of shape (m, Tx, 10)

    # When you pass a tensor through a Dense layer, it multiplies the values in_
    ↪the last dimension of the input tensor with
    # the layer's weights and adds the bias, followed by applying the_
    ↪activation function, resulting in a new last dimension as
    # specified by the number of units in the Dense layer.

    # Use 'densor2' to propagate e through a small fully-connected neural_
    ↪network to compute the "energies" variable energies.
    energies = densor2(e)
    # 'energies' is of shape (m, Tx, 1)

    # Use "activator" on "energies" to compute the attention weights "alphas"
    alphas = activator(energies)
    # 'alphas' is of shape (m,Tx,1)
    # 'alphas' matrix values change when context<t> changes (as timestep t_
    ↪changes).
```

```

    # Use dotor together with "alphas" and "a", to compute the context vector
    → to be given to the next (post-attention) LSTM-cell
    context = dotor([alphas,a])
    # For every 2*n_a vector in the 'a', there is one value in 'alphas' which
    → is along the T_x direction.
    # This value will be multiplied throughout all values of the 2*n_a vector.
    → Basically we are scaling the 2*n_a vector.
    # The scaling process is done for each time step and the vectors along the
    → T_x direction are added.
    # The shape of 'context' is (m, 2*n_a)

    # Another alternative method that we could have done
    # context = np.sum(alphas*a, axis = 1)

    return context

```

```

[47]: # UNIT TEST
def one_step_attention_test(target):

    m = 10
    Tx = 30
    n_a = 32
    n_s = 64
    #np.random.seed(10)
    a = np.random.uniform(1, 0, (m, Tx, 2 * n_a)).astype(np.float32)
    s_prev = np.random.uniform(1, 0, (m, n_s)).astype(np.float32) * 1
    context = target(a, s_prev)

    assert type(context) == tf.python.framework.ops.EagerTensor, "Unexpected
    → type. It should be a Tensor"
    assert tuple(context.shape) == (m, 1, n_s), "Unexpected output shape"
    assert np.all(context.numpy() > 0), "All output values must be > 0 in this
    → example"
    assert np.all(context.numpy() < 1), "All output values must be < 1 in this
    → example"

    #assert np.allclose(context[0][0][0:5].numpy(), [0.50877404, 0.57160693, 0.
    → 45448175, 0.50074816, 0.53651875]), "Unexpected values in the result"
    print("\033[92mAll tests passed!")

one_step_attention_test(one_step_attention)

```

All tests passed!

```
[48]: n_a = 32 # number of units for the pre-attention, bi-directional LSTM's hidden
      ↪ state 'a'
      n_s = 64 # number of units for the post-attention LSTM's hidden state "s"

      # This is the post attention LSTM cell.
      post_activation_LSTM_cell = LSTM(n_s, return_state = True) # Please do not
      ↪ modify this global variable.
      # In a standard LSTM, return_state=True will return the last hidden state
      ↪ (a<Tx>) and the last cell state (c<Tx>), along with
      # the output sequence (if return_sequences=True) ([a<1>,a<2>,...a<Tx>]) or the
      ↪ last output (a<Tx>) (if return_sequences=False).

      # In a standard LSTM, if return_state is False, then it will not return the
      ↪ last cell state (c<Tx>) and will only return
      # the last hidden state a<Tx>.

      # In a Bidirectional LSTM, return_sequences=True ensures that you get the
      ↪ output (hidden states) from both directions for
      # each time step of the input sequence.

      output_layer = Dense(len(machine_vocab), activation=softmax)
```

```
[49]: def modelf(Tx, Ty, n_a, n_s, human_vocab_size, machine_vocab_size):
      """
      Arguments:
      Tx -- length of the input sequence
      Ty -- length of the output sequence
      n_a -- hidden state size of the Bi-LSTM
      n_s -- hidden state size of the post-attention LSTM
      human_vocab_size -- size of the python dictionary "human_vocab"
      machine_vocab_size -- size of the python dictionary "machine_vocab"

      Returns:
      model -- Keras model instance
      """

      # Define the inputs of your model with a shape (Tx, human_vocab_size)
      X = Input(shape=(Tx, human_vocab_size))

      # Define s0 (initial hidden state) and c0 (initial cell state)
      # for the decoder LSTM with shape (n_s,)
      s0 = Input(shape=(n_s,), name='s0')
      c0 = Input(shape=(n_s,), name='c0')
      # Note the 'comma' is required in 'shape = (n_s,)' because in Keras when we
      ↪ define an input shape, we have to pass a tuple
      # representing the dimensions of the input.
```

```

# 's0' and 'c0' are just for 1 timestep (the initial one)

# hidden state
s = s0
# cell state
c = c0

# Initialize empty list of outputs
outputs = []

# Define the pre-attention Bi-LSTM
a = Bidirectional(LSTM(n_a, return_sequences = True),
merge_mode='concat')(X)
# In a Bidirectional LSTM, return_sequences=True ensures that you get the
output (hidden states) from both directions for
# each time step of the input sequence.
# LSTM, Keras automatically initializes the hidden state and the cell state
to zero vectors by default if we don't
# specify them. And automatically  $c_{<1>}$ ,  $a_{<1>}$  are used to calculate  $a_{<2>}$ .

# Iterate for Ty steps
for t in range(Ty):

    # Perform one step of the attention mechanism to get back the context
vector at step t
    context = one_step_attention(a, s)

    # Apply the post-attention LSTM cell to the "context" vector.
    _, s, c = post_activation_LSTM_cell(context, initial_state = [s,c])

    # Note - The below line of code doesn't work because when we do it
iteratively in keras we have to explicitly mention
    # the inputs (_, s, c = post_activation_LSTM_cell(initial_state =
[s,c])(context))

    # Apply Dense layer to the hidden state output of the post-attention
LSTM
    out = output_layer(s)

    # Append "out" to the "outputs" list
    outputs.append(out)

# Create model instance taking three inputs and returning the list of
outputs
model = Model(inputs=[X,s0,c0], outputs=outputs)

```



```

return model

# Note 1
# The pre-attention Bidirectional LSTM is called the 'encoder'. Usually the
  ↳ 'encoder' calculates the output all at once. This is
# because the output of one time step is not fed as input to calculate the
  ↳ output of next time step.
# Whereas the 'post_activation_LSTM_cell' is called the 'decoder'. The
  ↳ 'context' of each time step is fed as input to calculate
# the output of next time step. Hence we calculate the outputs of the time
  ↳ steps iteratively.

# Note 2
# We could have defined 's0' and 'c0' to be zero vectors of respective sizes
  ↳ (n_s) and
# could have defined 'model = Model(inputs=[X], outputs=outputs)'.
# But if 's0' and 'c0' are hardcoded, it would decrease the flexibility and
  ↳ they would always be zero vectors and
# can't be input.

```

```

[50]: # UNIT TEST
from test_utils import *

def model_test(target):
    Tx = 30
    n_a = 32
    n_s = 64
    len_human_vocab = 37
    len_machine_vocab = 11

    model = target(Tx, Ty, n_a, n_s, len_human_vocab, len_machine_vocab)

    print(summary(model))

    expected_summary = [['InputLayer', [(None, 30, 37)], 0],
                        ['InputLayer', [(None, 64)], 0],
                        ['Bidirectional', (None, 30, 64), 17920],
                        ['RepeatVector', (None, 30, 64), 0, 30],
                        ['Concatenate', (None, 30, 128), 0],
                        ['Dense', (None, 30, 10), 1290, 'tanh'],
                        ['Dense', (None, 30, 1), 11, 'relu'],
                        ['Activation', (None, 30, 1), 0],
                        ['Dot', (None, 1, 64), 0],
                        ['InputLayer', [(None, 64)], 0],

```

```

        ['LSTM',[(None, 64), (None, 64), (None, 64)],
↪33024,[(None, 1, 64), (None, 64), (None, 64)],'tanh'],
        ['Dense', (None, 11), 715, 'softmax']]

    assert len(model.outputs) == 10, f"Wrong output shape. Expected 10 !=
↪{len(model.outputs)}"

    comparator(summary(model), expected_summary)

model_test(model)

```

```

[['InputLayer', [(None, 30, 37)], 0], ['InputLayer', [(None, 64)], 0],
['Bidirectional', (None, 30, 64), 17920], ['RepeatVector', (None, 30, 64), 0,
30], ['Concatenate', (None, 30, 128), 0], ['Dense', (None, 30, 10), 1290,
'tanh'], ['Dense', (None, 30, 1), 11, 'relu'], ['Activation', (None, 30, 1), 0],
['Dot', (None, 1, 64), 0], ['InputLayer', [(None, 64)], 0], ['LSTM', [(None,
64), (None, 64), (None, 64)], 33024, [(None, 1, 64), (None, 64), (None, 64)],
'tanh'], ['Dense', (None, 11), 715, 'softmax']]
All tests passed!

```

```
[51]: model = model(Tx, Ty, n_a, n_s, len(human_vocab), len(machine_vocab))
```

```
[ ]: model.summary()
```

```
[53]: # Compiling the Model
```

```
[54]: opt = Adam(lr=0.005, beta_1=0.9, beta_2=0.999, decay=0.01) # Adam(...)
model.compile(loss = 'categorical_crossentropy', optimizer = opt, metrics =
↪['accuracy'])
```

```
[55]: # UNIT TESTS
assert opt.lr == 0.005, "Set the lr parameter to 0.005"
assert opt.beta_1 == 0.9, "Set the beta_1 parameter to 0.9"
assert opt.beta_2 == 0.999, "Set the beta_2 parameter to 0.999"
assert opt.decay == 0.01, "Set the decay parameter to 0.01"
assert model.loss == "categorical_crossentropy", "Wrong loss. Use
↪'categorical_crossentropy'"
assert model.optimizer == opt, "Use the optimizer that you have instantiated"
assert model.compiled_metrics._user_metrics[0] == 'accuracy', "set metrics to
↪['accuracy']"

print("\033[92mAll tests passed!")

```

All tests passed!

```
[56]: # Model Fitting
```

```
[57]: s0 = np.zeros((m, n_s))
      c0 = np.zeros((m, n_s))

      # The 'outputs' to be fed to the model is a list of Ty elements where each
      #   ↳ element is of shape (m, len(machine_vocab)).
      # But 'Yoh' is of shape (m, Ty, len(machine_vocab))
      # So we swap the axes 0 and 1 to make 'Yoh' of the shape
      #   ↳ (Ty, m, len(machine_vocab)).
      # Then we do 'list(swapped_Yoh)' so that 'Yoh' becomes a list consisting of Ty
      #   ↳ elements where
      # each element of size (m, Ty, len(machine_vocab))

      outputs = list(Yoh.swapaxes(0,1))
```

```
[58]: model.fit([Xoh, s0, c0], outputs, epochs=1, batch_size=100)
```

```
100/100 [=====] - 9s 86ms/step - loss: 16.5680 -
dense_5_loss: 1.1589 - dense_5_1_loss: 1.0174 - dense_5_2_loss: 1.8156 -
dense_5_3_loss: 2.6593 - dense_5_4_loss: 0.7361 - dense_5_5_loss: 1.2673 -
dense_5_6_loss: 2.6730 - dense_5_7_loss: 0.9684 - dense_5_8_loss: 1.7123 -
dense_5_9_loss: 2.5598 - dense_5_accuracy: 0.5120 - dense_5_1_accuracy: 0.6828 -
dense_5_2_accuracy: 0.2907 - dense_5_3_accuracy: 0.0836 - dense_5_4_accuracy:
0.9513 - dense_5_5_accuracy: 0.3205 - dense_5_6_accuracy: 0.0500 -
dense_5_7_accuracy: 0.9391 - dense_5_8_accuracy: 0.2542 - dense_5_9_accuracy:
0.1061
```

```
[58]: <tensorflow.python.keras.callbacks.History at 0x7fa59e17cb50>
```

```
[59]: # The below model has been run for longer time and the weights have been saved

      model.load_weights('models/model.h5')
```

```
[60]: # Results with the saved model weights

      EXAMPLES = ['3 May 1979', '5 April 09', '21th of August 2016', 'Tue 10 Jul
      #   ↳ 2007', 'Saturday May 9 2018', 'March 3 2001', 'March 3rd 2001', '1 March
      #   ↳ 2001']

      s00 = np.zeros((1, n_s))
      c00 = np.zeros((1, n_s))
      for example in EXAMPLES:
          source = string_to_int(example, Tx, human_vocab)
          #print(source)
          source = np.array(list(map(lambda x: to_categorical(x,
          #   ↳ num_classes=len(human_vocab)), source))).swapaxes(0,1)
          source = np.swapaxes(source, 0, 1)
          source = np.expand_dims(source, axis=0)
          prediction = model.predict([source, s00, c00])
```

```

prediction = np.argmax(prediction, axis = -1)
output = [inv_machine_vocab[int(i)] for i in prediction]
print("source:", example)
print("output:", ''.join(output), "\n")

```

source: 3 May 1979
output: 1979-05-33

source: 5 April 09
output: 2009-04-05

source: 21th of August 2016
output: 2016-08-20

source: Tue 10 Jul 2007
output: 2007-07-10

source: Saturday May 9 2018
output: 2018-05-09

source: March 3 2001
output: 2001-03-03

source: March 3rd 2001
output: 2001-03-03

source: 1 March 2001
output: 2001-03-01

[61]: *# Note*
Machine translation models can be used to map from one sequence to another.
They are useful not just for translating human languages (like
→French→English) but also for tasks like
date format translation.

A network using an attention mechanism can translate from inputs of length T_x
→to outputs of length T_y , where
T_x and T_y can be different.

The model that we built here can be used to translate from one language to
→another, such as translating from English to Hindi.
However, language translation requires massive datasets and usually takes
→days of training on GPUs.